MAGS: Learning LLMS

By: Pulkith, Pragya, Stanley, Caleb, Shailesh

Humans vs LLMs

- Lack of Agency
  - Figure, Tesla, etc...
- **Lack of Memory, Learning, and World Model**
  - **Our Goal**
- Lack of True Understanding, Reasoning, Novelty
  - ???

Motivation & Problem

Memory and Learning

- LLMs cannot **learn and remember**
  - Test-time emulation with context windows is not viable
    - Quadratic scaling
    - Attention decay and ephemeral context
    - Stateless session
  - Retrieval Augmented Generation (RAG) can add information
    - But rigid, one way static transfer of knowledge
    - Cannot reconcile new or conflicting information permanently

Motivation & Problem

Current SoTA

- CAMELoT, MIT+IBM
  - Added Memory Blocks To Remember Information
  - Fixed Memory Size, Overwrites Past Memory, No relationships between memory, No reconciliation
- Continual Learning
  - Catastrophic Forgetting, Expensive, Very Slow, not optimal
  - Google Titans (RMTs): Updates Attention loses fine details, and slow + 2M token max
- AriGraph
  - Adding Dynamic Graph-Based RAG
  - LLM has no representation of memory, only adds information (no updating+consolidating), no importance of different episodes, memory not connection across episodes

Motivation & Problem

Technical Approach

M	A	G	S
Memory	**Augmented**	**Graph**	**Scaling**
Episodic and Semantic Memory Blocks (RAG) along with Long-term vs Working Memory Supports separation of general knowledge, experience, and importance.	Compatible with any vanilla frozen LLM. Fine-tuned for engram+recall steps, so LLM can interact with memory. LLM can query knowledge before answering and store novel experiences and learnings.	Liquid Knowledge Graphs, that can add+remove nodes, change edge weights, consolidate and forget knowledge. Replicates human Hebbian Plasticity, Neurogenesis, and Synaptic Pruning.	Allows for context scaling limited by memory constraints rather than attention decay. Similar time inference. Chat History is part of Working memory, so agents are more modular.

Technical Approach

Structure

Knowledge broken down into Long-Term (important info), and Working Memory (chat history). Memory is moved from Working to Long-Term based on usage, ‘suprise’, or importance.

Recall

Anchor Nodes selected from query, BFS over decaying edge weights. SCCs consolidated. Can also query for connection between nodes.

Engram

Updates nodes and edge weights based on query importance. Can add or update information.

Validation

Custom game with nonsensical rules (so mode cannot reason). LLM can store newfound experiences, rule changes, surprising things, general game rules, etc...

Mathematical Formulation - GRPO

Let $\pi_\phi(a|s)$ be the policy with parameters ϕ and $\pi_{\phi_{\text{old}}}(a|s)$ be the policy before the update. Define the probability ratio:

$$r_t(\phi) = \frac{\pi_\phi(a_t|s_t)}{\pi_{\phi_{\text{old}}}(a_t|s_t)}.$$

Let A_t be the advantage estimate at time t and define the group-relative advantage as:

$$\hat{A}_t^{\text{GRPO}} = A_t - \frac{1}{|\mathcal{G}(t)|} \sum_{t' \in \mathcal{G}(t)} A_{t'},$$

where $\mathcal{G}(t)$ is the set of experiences in the group corresponding to time t.

Then the GRPO objective is:

$$L^{\text{GRPO}}(\phi) = \mathbb{E}_t \left[\min \left(r_t(\phi)\, \hat{A}_t^{\text{GRPO}},\ \text{clip}\left(r_t(\phi),\, 1-\epsilon,\, 1+\epsilon\right)\, \hat{A}_t^{\text{GRPO}} \right) \right],$$

where ϵ is a hyperparameter that limits the extent of policy updates.

Mathematical Formulation - MAGS

$$\min_{\theta,\phi,\psi} \mathcal{L} = \mathbb{E}_{(x,y)\sim\mathcal{D}} \left[\ell\Big(f_\theta\Big(x, g_\phi\big(M(x;\psi)\big)\Big), y \Big) \right] - \lambda\, \mathbb{E}_{x\sim\mathcal{D}} \left[R\Big(M(x;\psi), g_\phi\big(M(x;\psi)\big) \Big) \right] + \mu\, \mathcal{L}_{\text{graph}}\Big(M(x;\psi)\Big) + \nu\, \mathcal{R}(\theta,\phi,\psi).$$

We define our objective as a composite loss that jointly optimizes the language model's prediction accuracy, memory retrieval/reinforcement, and dynamic graph structure regularization. Let:

- θ denote the parameters of the underlying LLM.
- ϕ denote the parameters governing the memory retrieval and update module.
- ψ denote the parameters controlling the dynamic graph (i.e., Liquid Knowledge Graph) structure.
- $\mathcal{D} = \{(x, y)\}$ be the dataset of input–output pairs.
- $M(x;\psi)$ be the memory representation extracted from input x (including both episodic and semantic components).
- $g_\phi\big(M(x;\psi)\big)$ be the memory retrieval function that selects relevant memory nodes.
- $f_\theta(\cdot)$ be the generative function of the LLM augmented with the retrieved memory.
- $\ell(\cdot,\cdot)$ be a standard prediction loss (e.g., cross-entropy).

1. **Prediction Loss Term:**

$$\mathbb{E}_{(x,y)\sim\mathcal{D}} \left[\ell\Big(f_\theta\Big(x, g_\phi\big(M(x;\psi)\big)\Big), y \Big) \right]$$

ensures that the model's predictions are accurate given the input and the augmented memory.

2. **Memory Reward Term:**

$$-\lambda\, \mathbb{E}_{x\sim\mathcal{D}} \left[R\Big(M(x;\psi), g_\phi\big(M(x;\psi)\big) \Big) \right]$$

where $R(\cdot)$ is a reinforcement signal (e.g., derived from Group Relative Policy Optimization) that rewards effective memory recall and engram updates. The hyperparameter λ balances its influence.

3. **Graph Regularization Term:**

$$\mu\, \mathcal{L}_{\text{graph}}\Big(M(x;\psi)\Big)$$

is a penalty term (which may include terms for edge density, conflict resolution, and pruning cost) to maintain an efficient, sparse, and interpretable Liquid Knowledge Graph. The hyperparameter μ regulates its strength.

4. **Regularization Term:**

$$\nu\, \mathcal{R}(\theta,\phi,\psi)$$

is a composite regularization term (including, for example, L_2 norms, memory capacity constraints, and complexity penalties) that ensures the overall system remains computationally feasible and stable. The hyperparameter ν controls its weight.

Current Results

34%, 17%

Improvements in accuracy of Base Model, Base Mode + Static RAG on 100-MCQ question answering (2K context window). 83% total.

119, 491

Average Node Count, Edge Count after answering 100 questions. Typically, <3 nodes pruned. (No consolidation as of yet)

~60%

Instruction following with PPO (using meta tokens correctly in the right contexts)

~50%

Correct conflict resolution with new and conflicting information.

Current Issues

Models

The model is either too conservative or liberal with memory allocation. Doesn't always follow instructions (small model so we can fine-tune)

Data

Difficult to get high-fidelity information for RL. Currently, multishot prompting ChatGPT, but not much uniqueness or variety in this.

Compute

Takes 12–36 hours to test most architectural changes, so difficult to have iteration cycles.

Tuning

Difficult to find (and test) optimal parameters, ~100 parameters like edge decay, query document count, max density, etc..

Next Steps

Different Models

Try better prompting techniques, and slightly larger models to see if instruction following and expected behaviors are better.

Label more Data

Label more diverse and variety data so that the RL algorithm has better information and data to train on.

Cloud Computing

Try using virtual GPUs (and write the code to use cloud GPUs) to speed up the training and iteration process.

Parallelize + Finalize

Test out a lot of different configurations at the same time on length tasks to find some optimal candidates to pursue. Also add remaining features like consolidation, memory movement, etc..

PITCH DECK

3D MODERN BACKGROUND

HOW TO USE THIS PRESENTATION

Canva

Click on the "Canva" button under this presentation preview. Start editing your presentation. You need to sign in to your Canva account.

Powerpoint

Export this design from Canva as a PowerPoint template. Open the design in Canva. This will provide you with all the fonts used and elements used in this presentation as listed on page 21. Learn more on slide 3.

Google Slides

Export this design from Canva as a GoogleSlide template. This will provide you with all the fonts used and elements used in this presentation as listed on page 21. Learn more on slide 4.

TABLE OF CONTENTS

01.

ABOUT US

Elaborate on what you want to discuss.

HELLO THERE!

I'm Raina, and I'll be sharing with you my beautiful ideas. Follow me a @reallygreatsite to learn more.

ABOUT US

What we do?

Briefly elaborate on what you want to discuss.

Who we are?

Briefly elaborate on what you want to discuss.

ADD A TABLE

Add a main point What do you do during your free time? Write them here!	**Add a main point** What do you do during your free time? Write them here!	**Add a main point** What do you do during your free time? Write them here!
Add a main point What do you do during your free time? Write them here!	**Add a main point** What do you do during your free time? Write them here!	**Add a main point** What do you do during your free time? Write them here!

WRITE AN ORIGINAL STATEMENT OR INSPIRING QUOTE

— INCLUDE A CREDIT, CITATION, OR SUPPORTING MESSAGE

ADD A TEAM MEMBERS PAGE

Name

Title or Position

Name

Title or Position

Name

Title or Position

02.

OUR PROJECTS

Elaborate on what you want to discuss.

YOUR PROJECT

Presentations are communication tools that can be used as demonstrations, lectures, speeches, reports, and more. It is mostly presented before an audience. It serves a variety of purposes, making presentations powerful tools for convincing and teaching.

Year: 2022

Role: Write your role

S	W	O	T
STRENGTHS	**WEAKNESSES**	**OPPORTUNITIES**	**THREATS**
What are you doing well? What sets you apart? What are your good qualities?	Where do you need to improve? Are resources adequate? What do others do better than you?	What are your goals? Are demands shifting? How can it be improved?	What are the blockers you're facing? What are factors outside of your control?

OUR SERVICES

Service One

Elaborate on what you want to discuss.

Service Two

Elaborate on what you want to discuss.

Service Three

Elaborate on what you want to discuss.

Service Four

Elaborate on what you want to discuss.

WRITE YOUR BIG TOPIC OR IDEA

Presentations are communication tools that can be used as demonstrations, lectures, speeches, reports, and more.

It serves a variety of purposes, making presentations powerful tools for convincing and teaching.

03.

STATS & NUMBERS

Elaborate on what you want to discuss.

2 OUT OF 5

Elaborate on the featured statistic.

37%

Elaborate on the featured statistic.

95%

Elaborate on the featured statistic.

12 MILLION

Elaborate on the featured statistic.

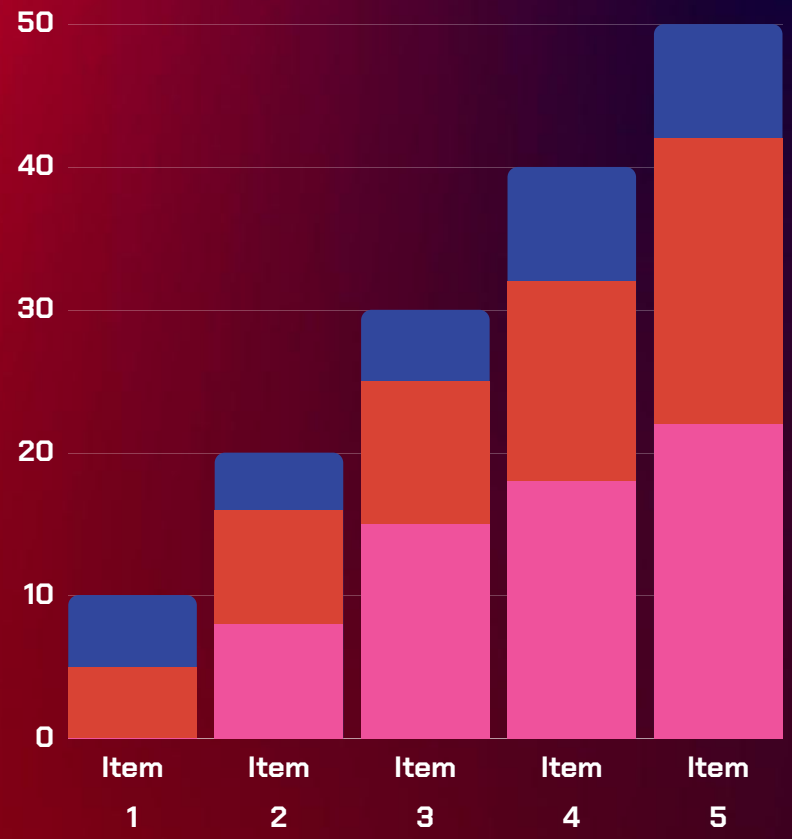

WRITE YOUR TOPIC OR IDEA

Elaborate on what you want to discuss.

INFOGRAPHIC SLIDE REPORT

Elaborate on the featured statistic

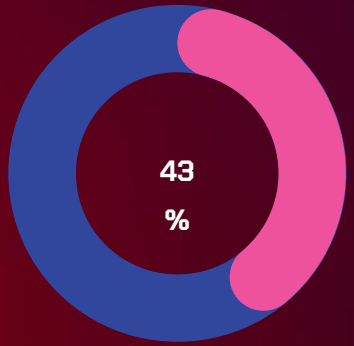

Elaborate on the featured statistic.

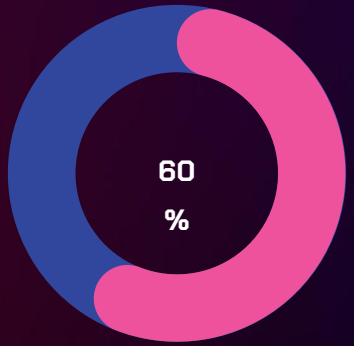

Elaborate on the featured statistic.

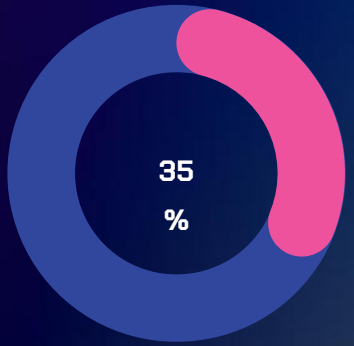

Elaborate on the featured statistic.

ADD A TIMELINE PAGE

Add a main point

Elaborate on a personal detail you want to share.

Add a main point

Elaborate on a personal detail you want to share.

Add a main point

Elaborate on a personal detail you want to share.

CONTACT US

123 Anywhere St., Any City,ST 12345	
123-456-7890	
hello@reallygreatsite.com	
@reallygreatsite	
reallygreatsite.com	

THANK YOU!

Write a closing statement or call-to-action here.

CREDITS

This presentation template is free for everyone to use thanks to the following:

SlidesCarnival for the presentation template
Pexels for the photos

Happy designing!

FONTS IN THIS PRESENTATION

This presentation template uses the following free fonts:

Titles: Chakra Petch
Headers: Chakra Petch
Body Copy: Montserrat

You can find these fonts online too.

Happy designing!

RESOURCE PAGE

Use these design resources in your Canva Presentation.
Happy designing!

Don't forget to delete this page before presenting.

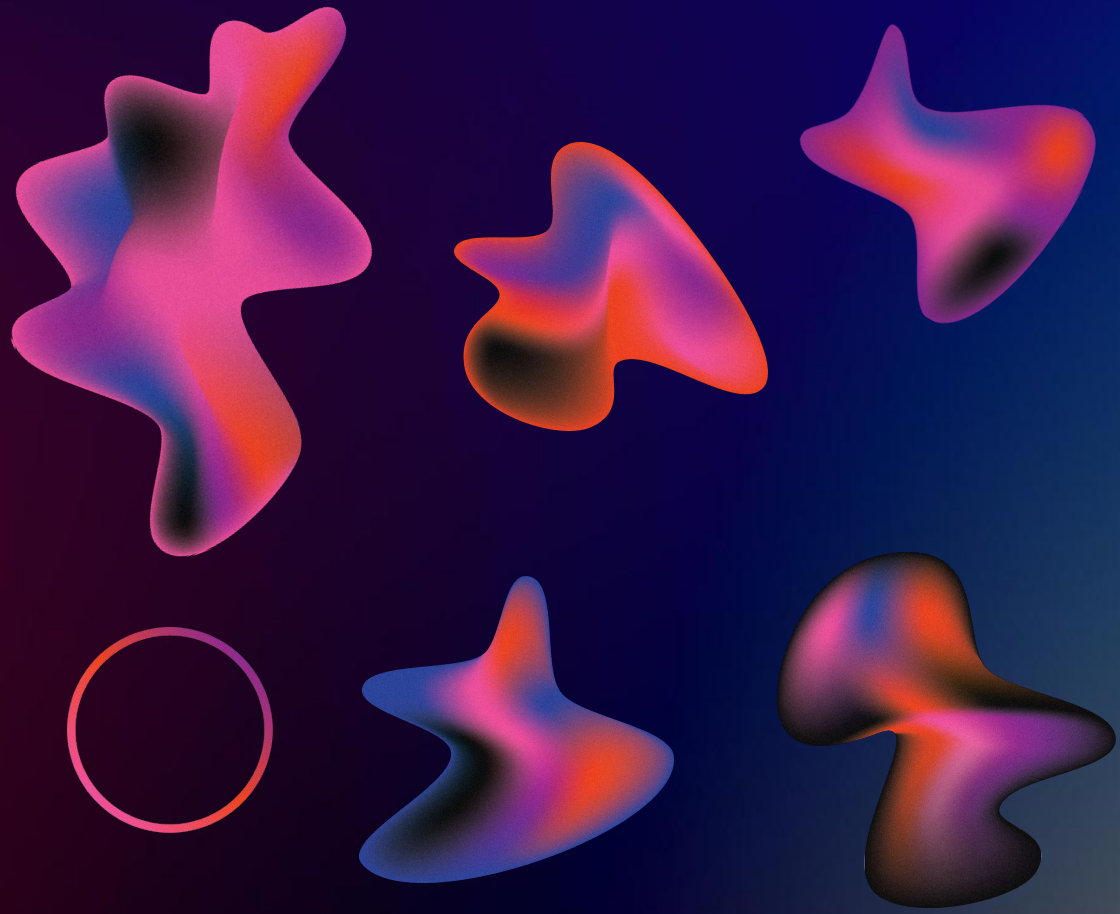

FREE TEMPLATES FOR ALL YOUR PRESENTATION NEEDS

For PowerPoint, Google Slides, and Canva

100% free for personal or commercial use

Ready to use, professional, and customizable

Blow your audience away with attractive visuals